

QWK by layer — llama-embed-nemotron-8b · MLP Regression · Out-of-Distribution
one line per dataset | dashed black = XGBoost baseline

